

CONTACT INFORMATION	School of Computer Science, Shanghai Jiao Tong University, 800 Dongchuan Road, Shanghai, China Homepage: https://zhitengli.github.io/	ieezhitengli@gmail.com Tel: +86-134-140-18498 Google Scholar, Github
RESEARCH INTERESTS	LLM pretraining and scaling; Image/Video generation; Model compression and inference acceleration (e.g., Quantization); Data synthesis.	
EDUCATION	Shanghai Jiao Tong University , Shanghai, China M.S. Student, Computer Science and Technology • GPA: 3.94/4.0 (Top 5%) Shanghai Jiao Tong University , Shanghai, China B.S., Computer Science and Technology (IEEE Honor Class) • GPA: 4.0/4.3 (Rank 1/113)	Sep 2023 – Mar 2026 Advisor: <i>Prof. Yulun Zhang</i> Sep 2019 – Jun 2023
PROFESSIONAL EXPERIENCE	ByteDance Seed, Beijing, China Research Intern Project: LLM pretraining with MuP & Stable Scaling Law. • MuP : Designed MuP for Muon, Weight Normalization, and Seed-LLM family. • Significantly stabilized training with no loss spikes . • Perfect hyperparameter transfer across 12× width and depth scaling. • Significant and consistent performance gain (5%-10%) on benchmark. • Stable Scaling Law : Proposed a novel scaling paradigm that decouples model scaling & data scaling, with a dual scaling scheme for learning rate & weight decay. • Maintaining lossless performance at an extreme batch size up to 128M . • Validated on a series of scaled models up to 2T parameters and 35T tokens , with better scaling trends across model and data scales. • Only 1/3 compute overhead for HPs search versus traditional scaling law. • Core contributor to Seed 2.0, Seed 2.5, and Seed 3.0. Xiaohongshu, Shanghai, China Research Intern Project: Video DiT acceleration (applied our QuantCache to real-world business scenarios). Sony AI, Remote Research Intern Project: Generative Data Augmentation (GenDataAgent is accepted at ICLR 2025). Amazon Shanghai AI Lab (ASAILab), Shanghai, China Applied Scientist Intern Project: Contributed to DGL & Graph Transformer research.	May 2025 – Apr 2026 Mentors: <i>Huizhuo Yuan, Prof. Quanquan Gu</i> Jan 2025 – Apr 2025 Mentor: <i>Shuang Sun</i> Dec 2023 – May 2024 Mentor: <i>Lele Chen</i> Jun 2022 – Sep 2023 Mentors: <i>Mufei Li, Minjie Wang</i>

PUBLICATIONS 5×ICLR, 1×ICCV, 1×CVPR, 5×Preprints. * indicates equal contribution.

- [CVPR 2026] **AdaSVD: Adaptive Singular Value Decomposition for Large Language Models**.
Zhiteng Li*, Mingyuan Xia*, Jingyuan Zhang, Zheng Hui, Haotong Qin, Linghe Kong, Yulun Zhang, and Xiaokang Yang.
- [ICLR 2026] **DVD-Quant: Data-free Video Diffusion Transformers Quantization**.
Zhiteng Li*, Hanxuan Li*, Junyi Wu, Kai Liu, Haotong Qin, Linghe Kong, Guihai Chen, Yulun Zhang, and Xiaokang Yang.
- [ICLR 2026] **Quant-dLLM: Post-Training Extreme Low-Bit Quantization for Diffusion Large Language Models**.
Tianao Zhang*, Zhiteng Li*, Xianglong Yan, Haotong Qin, Yong Guo, and Yulun Zhang.

4. [ICCV 2025] **QuantCache: Adaptive Importance-Guided Quantization with Hierarchical Latent and Layer Caching for Video Generation.**
Junyi Wu*, **Zhiteng Li***, Zheng Hui, Linghe Kong, Yulun Zhang, and Xiaokang Yang.
5. [ICLR 2025] **ARB-LLM: Alternating Refined Binarizations for Large Language Models.**
Zhiteng Li*, Xianglong Yan*, Tianao Zhang, Haotong Qin, Dong Xie, Jiang Tian, Zhongchao Shi, Linghe Kong, Yulun Zhang, and Xiaokang Yang.
6. [ICLR 2025] **GenDataAgent: On-the-fly Dataset Augmentation with Synthetic Data.**
Zhiteng Li, Lele Chen, Jerone Andrews, Yunhao Ba, Yulun Zhang, and Alice Xiang.
7. [ICLR 2026] **PT²-LLM: Post-Training Ternarization for Large Language Models.**
Xianglong Yan*, Chengzhu Bao*, **Zhiteng Li**, Tianao Zhang, Kaicheng Yang, Haotong Qin, Ruobing Xie, Xingwu Sun, and Yulun Zhang.

PREPRINTS

1. **VEQ: Modality-Adaptive Quantization for MoE Vision-Language Models.**
Guangshuo Qin*, **Zhiteng Li***, Zheng Chen, Weihang Zhang, Linghe Kong, and Yulun Zhang.
2. **ReCalKV: Low-Rank KV Cache Compression via Head Reordering and Offline Calibration.**
Xianglong Yan*, **Zhiteng Li***, Tianao Zhang, Haotong Qin, Linghe Kong, Yulun Zhang, and Xiaokang Yang.
3. **BinaryHPE: 3D Human Pose and Shape Estimation via Binarization.**
Zhiteng Li, Yulun Zhang, Jing Lin, Haotong Qin, Jinjin Gu, Xin Yuan, Linghe Kong, and Xiaokang Yang.
4. **D²Quant: Accurate Low-bit Post-Training Weight Quantization for LLMs.**
Xianglong Yan*, Chengzhu Bao*, **Zhiteng Li**, Tianao Zhang, Haotong Qin, Shaoqiu Zhang, Ruobing Xie, Xingwu Sun, and Yulun Zhang.
5. **Progressive Binarization with Semi-Structured Pruning for LLMs.**
Xianglong Yan, Tianao Zhang, **Zhiteng Li**, Haotong Qin, and Yulun Zhang.

ACADEMIC SERVICE

Reviewer

- International Conference on Machine Learning (ICML) 2025, 2026
- Advances in Neural Information Processing Systems (NeurIPS) 2025
- International Conference on Learning Representations (ICLR) 2025, 2026
- International Conference on Computer Vision (ICCV) 2025
- Computer Vision and Pattern Recognition (CVPR) 2025, 2026

HONERS AND AWARDS

- Yang Yuanqing Education Foundation 2025
- Excellent Graduate of Shanghai Jiao Tong University 2023
- Ruiyuan-Sequoia Talent Development Fund 2022
- National Scholarship for Undergraduate Excellence (Top 0.2% nationwide) 2021
- Excellent Undergraduate Scholarship 2020, 2021, 2022

SKILLS

- Programming: Python, Pytorch, C/C++, L^AT_EX